

Citation under Pressure: What Deployment Constraints Do to Legal Citation Integrity, and What Deterministic Verification Can Repair

Harsh DUBEY¹

New York University, New York, USA

Abstract. Sanctions over fabricated legal citations have made language-model drafting a problem for courts, and most research on the problem has produced detection benchmarks that score errors in finished briefs. This paper measures how the errors are produced. Six models from six vendors, from a 30B open-weight model to two frontier systems, drafted argument sections for 48 U.S. Supreme Court matters under five prompting conditions: baseline, a citation quota, a restriction to pre-1970 authority, a sanctions warning, and all three combined. Every citation in the 1,440 drafts was checked deterministically, with no language-model judge in any primary measurement: existence against an 18-million-row reporter index, quotations against a verbatim checker validated on seeded faults, and pinpoint pages against official reporter pagination from public archives. Under combined pressure the smallest model’s citation existence fell from 0.924 to 0.773, and verbatim quotation accuracy fell significantly for three of six models; only the strongest model was unaffected. In a second experiment the checker fed its findings back into revision. The strongest model reached 0.974 verbatim accuracy, and a per-quotation trace shows how: of its 89 flagged quotations, 4 were corrected and the rest were deleted or stripped of their quotation marks. False feedback produced the same deletions aimed at correct quotations, and five of six models ended less accurate than they started. We release the instruments, including a page-level pinpoint verifier built from public data.

Keywords. legal citation, hallucination, large language models, verification, evaluation

1. Introduction

Since 2023, courts have repeatedly sanctioned lawyers who filed briefs containing citations that a language model invented, beginning with *Mata v. Avianca* [1] and continuing in a growing docket that a public tracker now counts in the hundreds [2]. Most published work frames the problem as detection: given a brief that already contains errors, how many can a model or agent find [3,4]? Detection accuracy is the relevant measure for review of a filed brief. It does not tell a deployer which prompting conditions produce the errors, which is the question this paper addresses.

¹Corresponding Author: Harsh Dubey, New York University; harshdubeyjp@gmail.com.

This paper asks two production-side questions. First, which prompting conditions cause citation failure, and does the answer depend on the model? Prompts in practice carry constraints of the kind studied by Zhao et al. [5] for scholarly citations: a minimum count of authorities, a restriction on which authorities may be cited, and a warning about consequences. That study found that such constraints, temporal restrictions in particular, collapse citation existence even in capable models. We are not aware of a comparable study for law. Second, when a checker that can verify citations deterministically feeds its findings back into revision, does the model correct its drafts, fail to, or change the draft so that the checker’s tests pass without the quotations being corrected?

Both questions need an instrument the field has said does not exist in public: the best published detection agent misses nearly half of incorrect pinpoint citations (52.8 percent recall; open-weight detectors 19 to 51 percent), and its authors conclude that “no models can reliably detect wrong pincites” because page-level information lives behind commercial databases [3]. We built the instrument from public sources: citation existence and name pairing against a local index derived from CourtListener bulk data, verbatim quotation checking validated on seeded faults, and page-level pinpoint verification against the official reporter pagination in the Caselaw Access Project’s volume archives [6]. Every primary number reported here is produced by this instrument; no language-model judge is used in any primary measurement.

Our scope is the closed-book use of general-purpose models, the setting of most sanctioned filings documented to date [2,7]; retrieval-grounded commercial tools are a different setting, and Section 7 says what we expect to transfer.

The paper reports three results. First, pressure degrades citation integrity beginning with the smallest model, and the degradation is rarely accompanied by refusal; in the 1,152 drafts under the four non-baseline conditions (the pressured drafts) there were three refusals. Second, feedback from a deterministic checker leads capable models to delete flagged quotations far more often than to correct them, and false feedback leads them to delete correct ones. Third, in the strongest models the remaining errors are mostly genuine legal language attributed to the wrong case or the wrong page, and the released instruments detect both classes.

2. Related Work

Legal citation reliability. Liu et al. [3] build a detection benchmark from real appellate briefs with injected errors and a taxonomy grounded in sanctioned filings; we adopt their five error categories as outcome coding and measure the naturally occurring analogue their limitations section asks for. LegalCiteBench [4] shows that models asked closed-book for the citation of a named case almost never produce the right one and almost never abstain; our task asks the easier question of whether the cases a model chooses to cite exist, and the low abstention transfers. Dahl et al. [7] measured hallucinated authority in general-purpose chatbots, and Magesh et al. [8] in commercial legal research tools; together they establish the rate of hallucinated authority in deployed systems.

Constraints and fabrication. Zhao et al. [5] is the closest design: five prompting regimes over scholarly claims, verified by fuzzy bibliographic matching that left 36 to 61 percent of citations unresolved. We transplant their factorial to law with deterministic verification (0.2 percent unresolved) and a formal statistical model; their mechanism,

compliance with the temporal window filled with fabrications, transfers to law for some models and not others (Section 4).

Revision with execution feedback. The loop experiment takes its protocol from execution-feedback work in code generation [9], including a no-feedback control that separates feedback content from extra revision attempts. Our version adds a second control in which the feedback is false.

AI and law. Blair-Stanek and Van Durme [10] found at ICAIL that frontier models flip the winning party across identical zero-temperature runs; we ran each cell once, and Section 8 treats that instability as a limitation of this design. Evaluation methodology in this community has moved toward separating factual truth from reasoning soundness [11] and pairing automated metrics with the checks lawyers perform [12]; citation integrity is a check for which deterministic ground truth is available, and this paper automates it. Closest to our pinpoint instrument, Garneau et al. [13] found that retrieving the paragraph a proposition rests on is far harder than retrieving the case, using network and text signals to propose pincites; our verifier approaches the same gap from the other side, checking a generated pincite against the page text. Zhang et al. [14] show models miss overruled precedent, a failure adjacent to the wrong-authority errors we measure, and ConsRAG [15] mitigates hallucination through constrained retrieval, the strategy complementary to the post-hoc feedback our loop experiment tests. Two studies of model-drafted litigation documents frame the trust question: Trozze et al. [16] found mock jurors’ decisions unrelated to whether a complaint was written by ChatGPT or a lawyer, and Harasta et al. [17] found lawyers rate identical documents lower when told a model wrote them; machine verification of citation integrity offers a basis for trust that depends on neither perception. Gray et al. [18] generate legal arguments whose quality is judged by lawyers, and Pires et al. [19] evaluate legal writing automatically; neither checks the authority the text cites, which is the layer this paper adds. Citation-worthiness detection [20] asks where a citation is needed; the citation index behind our checker is the same kind of network studied in [21,22,23], used here to test whether a cited case exists.

3. Tasks, Conditions, and Instruments

Terms. A citation such as *Brown v. Board of Education*, 347 U.S. 483, 495 (1954) names the parties, the volume (347), the reporter (U.S., the official United States Reports), and the first page (483); the optional second page (495) is the *pincite*, which for a quotation must be the page where the quoted language appears. Throughout, *citation existence* is the share of adjudicable citations that resolve to a real case, and *strict quotation accuracy* is the share of scored quotations that reproduce the opinion verbatim. The five prompting conditions are called baseline, quota, temporal, stakes, and combined.

Matters. The matters are forty-eight Supreme Court cases from 1990 to 2020, drawn from a 500-case pool built from the Supreme Court Database [24]. Each packet holds the question presented, party names, a factual summary, and a lower-court opinion excerpt, with the Supreme Court’s decision excluded and every packet screened for outcome leakage. The model argues an assigned side (petitioner for even-numbered matters, respondent for odd) in an argument section of under 900 words; nothing in the design predicts outcomes, so the memorization confound of Supreme Court prediction tasks [25] is out of scope.

Conditions. The five conditions are baseline; quota (“cite at least eight distinct case authorities”); temporal (“every authority you cite MUST have been decided before 1970”); stakes (“this section will be filed with the court; citing authority that does not exist, or misquoting authority, risks sanctions under Rule 11”); and combined, which concatenates the three clauses. Each condition has one fixed template, run at temperature 0. The stakes clause names Rule 11 of the Federal Rules of Civil Procedure because it is the sanctions rule invoked in the sanctioned-filing cases [1]. Rule 11 does not govern Supreme Court filings, which fall under the Court’s own rules and inherent authority; the clause is a stakes signal, not a statement of the governing law. The conditions isolate mechanisms rather than simulate practice. A robustness check re-ran twelve matters under alternative phrasings of every condition.

Models. The models are Qwen3-30B-A3B, Mistral Small, Llama-4 Maverick, DeepSeek V4 Flash, GPT-5.4-mini, and Claude Sonnet 5, one per vendor, all accessed through one API gateway. The hypotheses concern how effects change with capability, so the set spans open-weight 30B to frontier with no vendor contributing two models; we do not assume a total ordering, and tables are ordered by baseline strict quotation accuracy. The gateway identifiers (qwen3-30b-a3b-instruct-2507, mistral-small-2603, llama-4-maverick, deepseek-v4-flash, gpt-5.4-mini, claude-sonnet-5) were run in August 2026 through OpenRouter with completion limits of 1,400 to 12,000 tokens. A seventh model, GLM-4.7-flash, completed 207 of 240 drafts and is reported in Section 4.

Adjudication. Citations are extracted with eyecite [26] and given one of three labels: exists, not found, or unverifiable (vendor citations and coverage gaps), so that index coverage is never counted as fabrication; 0.2 percent of the 12,353 citation records in the full run are unverifiable. Existence and name pairing resolve against an 18-million-row citation table. Quotations are located by balanced quotation marks and attributed to the case named in the same sentence, else to the nearest citation within 260 characters after the quotation, else to the nearest preceding one. Each is checked verbatim against the cited opinion’s text from a local-first chain of public sources; a quotation is *near miss* when every fragment reaches token coverage of at least 0.85 without matching exactly, and *inaccurate* otherwise. The strict rate for a draft d is

$$S_d = \frac{n_{\text{accurate}}}{n_{\text{accurate}} + n_{\text{near miss}} + n_{\text{inaccurate}}}, \quad (1)$$

and a lenient rate counting near misses as correct is reported alongside it. Before any measurement the quotation checker scored 100 of 100 on twenty synthetic drafts with planted genuine quotations, wrong attributions, single-word corruptions, fabrications, and formatting artifacts; because that is a self-authored test, Section 6 also audits its verdicts on real drafts. Pinpoint pages are verified against the official reporter pagination parsed from the Caselaw Access Project’s HTML volumes, which mark every page break; a quotation counts as on the cited page when its longest fragment lies within that page’s text, and adjacent-page matches are reported separately. Vendor citations (Westlaw and Lexis) are excluded from fabrication counts.

Human reference. The same instrument scored 482 clean, pre-ChatGPT federal appellate brief excerpts from the benchmark of Liu et al. [3]. Human lawyers reach 0.407 strict and 0.680 lenient quotation accuracy (578 scored quotations; the near misses are chiefly bracketed alterations and ellipses) and 0.931 citation existence, the shortfall be-




Figure 1. Strict quotation accuracy (left) and citation existence (right) by condition, one line per model. The dotted line is the human-lawyer reference under the identical instrument. The right panel’s axis starts at 0.7.

ing index coverage of state and unpublished reporters rather than fabrication. Human briefs cite more non-federal authority and quote statutes and records more often than the models do, so the human figures serve only as a reference for the severity of the verbatim standard.

Inference. The primary analysis is a citation-level logistic generalized estimating equation per model, $\text{logit} P(y_{imk} = 1) = \beta_0 + \beta_k$ for citation i in matter m under condition k , with exchangeable working correlation within matters [27], so that the 48 matters rather than the citations are the independent units. Each model has 8 contrasts (four conditions, two outcomes); we report Holm-corrected p values within model and call an effect significant only when it survives correction. Correction is within model, so statements that count models (“three of six”) are descriptive. Odds ratios, 95 percent confidence intervals, and raw and corrected p values for all 48 contrasts are in the released results. A matter-level cluster bootstrap over draft rates agreed in sign with every corrected result. For the revision experiment, which was pre-registered descriptively, we add matter-paired comparisons of final strict accuracy between arms (Wilcoxon signed-rank test and a bootstrap interval on the mean difference), chosen after the data were seen. Hypotheses and decision rules were frozen before the full run; every later analysis decision, including two findings withdrawn after our own verification passes, is recorded in a dated ledger released with the code.

4. Experiment 1: Pressure

Table 1 and Figure 1 report the rates.

Citation existence. Only the 30B model’s existence moves under pressure: from 0.924 at baseline to 0.773 under the combined condition, with the quota, temporal, and combined contrasts all significant after correction (combined OR 0.30, 95 percent CI 0.18 to 0.51, corrected $p < 0.001$). Mistral Small’s existence sits between 0.86 and 0.93 across conditions and falls under the temporal and combined conditions without reaching significance; the four other models stay between 0.93 and 1.00 in every condition.

Quotation accuracy. The temporal restriction is the condition that damages quoting. After correction, strict quotation accuracy falls significantly under temporal and

Table 1. Baseline versus the combined condition, rows ordered by baseline strict quotation accuracy. Quotation columns give the strict rate with the lenient rate in parentheses; existence columns give the rate with the raw count of citations not found in brackets. An asterisk marks a contrast significant after Holm correction within model.

Model	Quotes, base	Quotes, comb.	Exist, base	Exist, comb.	Temporal viol.
Qwen3-30B	0.175 (0.564)	0.072 (0.320)*	0.924 [28]	0.773 [90]*	20%
Mistral Small	0.207 (0.608)	0.136 (0.383)	0.901 [37]	0.868 [69]	35%
DeepSeek V4 Flash	0.281 (0.631)	0.120 (0.436)*	0.972 [10]	0.959 [20]	11%
Sonnet 5	0.333 (0.622)	0.427 (0.642)	0.991 [3]	0.998 [1]	0%
Llama-4 Maverick	0.402 (0.618)	0.179 (0.453)*	0.958 [9]	0.939 [31]	10%
GPT-5.4-mini	0.460 (0.722)	0.294 (0.571)	0.981 [8]	0.962 [27]	13%

combined pressure for Qwen3-30B (0.175 to 0.072) and Llama-4 (0.402 to 0.179), and under the combined condition for DeepSeek (0.281 to 0.120); GPT-5.4-mini and Mistral Small fall in the same direction without surviving correction, and Sonnet 5 does not fall at all. The Llama-4 row is the most consequential for practice: at baseline it quotes at the human reference level, 0.402 against 0.407, and under the temporal restriction its accuracy drops to 0.139 (OR 0.20, CI 0.10 to 0.40, corrected $p < 0.001$). The lenient rates in Table 1 move in the same direction as the strict rates for every model, so the result does not depend on how bracketed alterations are scored.

A pre-registered contrast that did not survive correction. Sonnet 5’s strict quotation accuracy is numerically highest under the combined condition (0.427 against 0.333 at baseline; uncorrected $p = 0.0292$, corrected $p = 0.233$), and its citations not found went from 3 of 315 at baseline to 1 of 449. Neither survives correction, the stakes condition alone moves neither outcome, and the temporal restriction shifts the cited cases toward older, better-memorized opinions (median decision year 1943 under the combined condition against 1986 at baseline), an untested alternative to any reading in terms of more conservative citing. We report it as a null result and draw no conclusion from it.

Compliance with the temporal rule. Violations of the pre-1970 restriction (Table 1) run from 35 percent of adjudicable citations for Mistral Small to 1 in 803 for Sonnet 5 under the temporal and combined conditions. Mistral Small violates the rule most while its existence and quotation figures sit mid-table, so compliance does not track the other outcomes. Zhao et al. [5] found near-zero violations with fabrication filling the window; here the smaller models both fabricate and violate.

Refusals and one model that stalls. Across the 1,152 pressured drafts there were three refusals (two from DeepSeek, one from GPT-5.4-mini). GLM-4.7-flash produced no output on 33 of its 240 drafts, 28 of them in the combined condition, reasoning until a 30,000-token budget was exhausted in two escalating attempts; because its missing drafts concentrate in the condition under comparison it is excluded from Table 1 and reported as a third failure mode, distinct from fabrication and from refusal.

5. Experiment 2: The Checker inside the Revision Loop

Combined-condition drafts entered a revision loop. Each round, the checker returned one line per failure (a citation that does not resolve; a quotation that does not appear verbatim in the cited opinion), the model revised, and the loop ended on a clean draft or

Table 2. Revision under true checker feedback, 24 episodes per model. Round-0 quotations, flagged (inaccurate or near miss) and accurate, traced to the final draft; strict accuracy round 0 to final; Clean counts episodes ending on a draft the checker passed.

Model	n	Flagged at round 0			Accurate at round 0		Strict	Clean
		corrected	dequoted	deleted	n	removed		
Qwen3-30B	170	0	8	30	16	0	0.099 to 0.107	0
Mistral Small	130	0	12	39	15	2	0.208 to 0.264	9
DeepSeek V4 Flash	183	0	67	101	28	12	0.181 to 0.645	21
Sonnet 5	89	4	13	72	68	17	0.471 to 0.974	23
Llama-4 Maverick	53	0	9	28	11	5	0.176 to 0.349	21
GPT-5.4-mini	37	2	5	21	28	2	0.360 to 0.719	22

Table 3. The two controls. Removed: accurate round-0 quotations (the n of Table 2) deleted, dequoted, or made inaccurate. Strict: final strict accuracy. True – control: matter-paired difference in final strict accuracy (mean, bootstrap 95 percent interval, Wilcoxon p), over matters with a scored quotation in both final drafts.

Model	False feedback				No feedback			
	removed	strict	true – false		removed	strict	true – none	
Qwen3-30B	4	0.093	0.01 [-0.01 to 0.04]	0.208	2	0.072	0.04 [0.00 to 0.09]	0.075
Mistral Small	8	0.122	0.12 [0.05 to 0.21]	0.012	2	0.198	0.03 [-0.01 to 0.08]	0.169
DeepSeek V4 Flash	24	0.153	0.39 [0.16 to 0.62]	0.016	11	0.152	0.49 [0.27 to 0.70]	0.005
Sonnet 5	56	0.266	0.69 [0.52 to 0.86]	0.0001	18	0.595	0.36 [0.23 to 0.51]	0.0003
Llama-4 Maverick	6	0.188	0.03 [-0.14 to 0.17]	0.713	8	0.109	0.29 [0.08 to 0.53]	0.039
GPT-5.4-mini	15	0.241	0.44 [0.23 to 0.66]	0.004	9	0.543	0.12 [-0.06 to 0.32]	0.293

after three revisions. A control with *false feedback* gave the same number of identically formatted lines pointed at items that were correct, and a *no-feedback* control, added after the pre-registered analysis, gave the same number of revision rounds with an instruction to revise as the model saw fit. The loop used the first 24 matters in the pre-registered order (twelve per side). Loop rates are means over drafts on 24 matters, so round-0 figures differ from the pooled citation-level rates over 48 matters in Table 1. Citation counts in this section include unverifiable citations. To see what revision did to each quotation, every scored quotation in a round-0 draft was traced to the final draft: *corrected* if the same or an edited quotation on the same citation is now accurate, *dequoted* if its words remain without quotation marks, *deleted* if neither.

Satisfying the checker by deletion. Under true feedback the final strict rate ranges from 0.107 for Qwen3-30B to 0.974 for Sonnet 5 (Table 2; the order is not monotone in the table’s baseline ordering), and the rate alone is consistent with repair. The trace shows that the rise comes from deletion. Sonnet 5 had 89 flagged quotations at round 0: 4 were corrected, 13 kept their words and lost their quotation marks, and 72 were deleted; 17 of its 68 accurate quotations were removed as well. Its scored quotations fell from 157 to 73 while its accurate ones went from 68 to 70 (21 newly added). DeepSeek and GPT-5.4-mini corrected 0 and 2 of theirs; Qwen3-30B corrected none and kept most of its flagged quotations in place. Citation existence, by contrast, was repaired: the loop existence rate rose from 0.739 to 0.822 for Qwen3-30B and from 0.930 to 1.000 for Llama-4 while citation counts fell by at most a tenth, so unresolvable citations were replaced with real ones rather than dropped. At every capability tested, checker feedback led models to




Figure 2. Strict quotation accuracy from round 0 to the final draft under true feedback, false feedback, and no feedback, one color per model. A rate can rise because accurate quotations were added or because inaccurate ones were removed; Tables 2 and 3 separate the two.

remove flagged quotations or convert them to unmarked paraphrase far more often than to correct them; neither outcome is flagged by the checker, and the revised draft does not indicate which claims were originally quoted.

The two controls. Table 3 and Figure 2 separate feedback content from revision itself. Without feedback, Sonnet 5’s strict rate still rose (0.471 to 0.595), and true feedback added 0.36 to the final rate over revision alone (interval 0.23 to 0.51, $p=0.0003$); the same holds for DeepSeek (0.49, $p=0.005$) and Llama-4 (0.29, $p=0.039$). For GPT-5.4-mini the difference between true feedback and no feedback is 0.12 with an interval that includes zero, so its apparent repair is not distinguishable from what revision alone does; for Mistral Small and Qwen3-30B neither arm moves far. Under false feedback the flagged items were correct, and the models treated them as they treated true flags: Sonnet 5 removed 56 of its 68 accurate quotations (18 under no feedback) and its strict rate fell from 0.471 to 0.266; GPT-5.4-mini removed 15 of 28 and fell to 0.241. Five of six models ended below their starting rate (Qwen3-30B by under a point); Llama-4, the exception, and Qwen3-30B are also the two that changed least under true feedback. The mechanism is deletion of correct quotations, not rewriting: inaccurate counts fell in every model under false feedback. Our checker’s own false-alarm rate, measured by the provenance audit of Section 6, was 1.3 percent of Sonnet 5’s inaccurate verdicts (8 of 617), low enough not to produce this effect; a detector with a false-positive rate in the tens of percent would be expected to produce it, and intermediate rates were not tested.

6. Residual Errors in the Strongest Models

Three error classes remain when fabrication is rare, each measured deterministically on all six models.

Paraphrase inside quotation marks. Decomposing every inaccurate quotation against the cited opinion’s own text, 55 percent of Sonnet 5’s inaccurate quotations (413 of them) are recognizable paraphrases of the correct case presented as verbatim quotation, 202 match nothing substantial in the cited opinion, and 131 cite authority outside the text store’s scope; GPT-5.4-mini’s split is 52 percent paraphrase. For the strongest models the most common inaccurate quotation is a paraphrase of the cited opinion presented inside quotation marks.

Real language, wrong case. Searching every inaccurate quotation against a segmented corpus of 8,392 Supreme Court opinions locates 463 quotations across the six models (114 for Sonnet 5, 66 for Qwen3-30B) whose language appears verbatim in a different opinion than the one cited, where the matched probe is at least six words (a ten-word window for long quotations) so that stock legal phrases do not count. Quotations found nowhere in the corpus (411 for Sonnet 5) include the paraphrases of the preceding paragraph and quotations from courts outside it. Because the attribution rule could itself misassign a quotation, the same search serves as an audit of the checker: it classified 8 of Sonnet 5’s inaccurate verdicts and 6 of Qwen3-30B’s as the checker’s own error, where the language was in the attributed case after all. A further screen for models reproducing the argued case’s own later-decided opinion from memory found 41 candidates, of which 32 were record or statute text present in the prompt and the rest doctrinal formulas that the argued opinions were themselves quoting; the finding was withdrawn and is recorded in the ledger.

Right case, wrong page. Where a pinpoint and a quotation co-occur on a U.S. Reports authority, the quotation lies on the cited page 77 percent of the time for Sonnet 5, between 42 and 64 percent for the other five, and on a different page of the cited case 16 to 45 percent of the time (adjacent-page matches make up the remainder); pins outside the cited case are rare at every tier. This measures, from public data, the error class on which published detection agents have the lowest recall [3].

Relevance judging: three failed calibrations. Whether a real, correctly quoted case supports the proposition it is cited for cannot be looked up. Three language-model panels calibrated against the expert-labeled content-misrepresentation items of Liu et al. [3], with a pre-declared bar of 0.75 accuracy on decisive verdicts, scored 0.500, 0.700 (misrepresentation recall 2 of 11), and 0.594 (a frontier panel with 15,000-character grounded excerpts; recall 4 of 17). No verdict from a failed panel was interpreted; the result joins evidence that agreement statistics overstate judge validity [28,29].

7. Discussion

Whether a model fabricates legal citations under realistic prompting depends on the model: for the 30B model a quota and an authority restriction triple the odds of a nonexistent citation, and the models above it hold existence but lose quotation accuracy, most under the temporal restriction. Restrictions on authority by date or jurisdiction are plausible drafting instructions.

Two practical conclusions follow from the revision experiment. Feeding checker output back into revision reduced the number of quotations and corrected few of them; a brief with fewer verbatim claims exposes fewer false ones, but a reader cannot tell which claims were demoted to paraphrase. The checker is therefore more useful as a gate on output, holding a draft until its citations resolve and its quotations match, than as feedback. Any system that feeds a probabilistic detector’s flags to a capable model should measure the detector’s false-positive rate first: the models that removed the most flagged quotations under true feedback also removed the most correct ones under false feedback.

Existence, verbatim quotation, and pinpoint page can be checked deterministically from public data, and the released verifier does so; court orders that require disclosure

of AI assistance [2] could additionally require the checker’s output on the filing. With retrieval-grounded tools we expect existence failures to shrink and the residual classes, paraphrase in quotation marks and wrong pages, to remain, since they arise in composition rather than recall.

8. Limitations

The study covers a single domain and task: Supreme Court merits argument, in one packet format, with one fixed template per condition and a phrasing check on twelve matters in place of a full paraphrase factorial. The task is closed-book, matching the sanctioned filings but not commercial legal tools, and the stakes clause names a sanctions rule that does not govern the forum. The strict verbatim standard penalizes bracketed alterations and ellipses for humans and models alike, and the human reference differs from the model drafts in citation mix. The provenance corpus covers 8,392 Supreme Court opinions and cannot adjudicate quotations from other courts. The false-feedback control is all false; intermediate false-positive rates were not tested. The no-feedback control and the per-quotation trace were added after the pre-registered analysis, and the tests in Section 5 were chosen after the data were seen. Each cell was generated once at temperature 0; given the run-to-run instability documented for legal tasks [10], the per-cell rates carry sampling noise that repeated runs would bound and this study does not. On a twelve-matter check under alternative phrasings of every condition (the four models available when it was designed), combined-condition existence stayed within 0.06 of the original phrasing and no existence result changed sign, but strict quotation rates moved by up to forty points in either direction (GPT-5.4-mini 0.528 to 0.125, Sonnet 5 0.376 to 0.684), so the quotation rates in Table 1 are specific to the templates used, and the within-model contrasts, which hold the template fixed, are the claims we rely on. Inference is per model at 48 matters, so effects of a few points are not detectable, and the six models are one per vendor rather than a sampled population, so capability, family, and training recipe are confounded.

9. Conclusion

Under prompting constraints of the kind found in practice, citation integrity degrades from the smallest model upward: nonexistent citations at 30B, misquotation through the middle of the range, and misattribution and wrong pages at the top. A deterministic checker built from public data detects each class. Fed back into revision, its output yields briefs with fewer quotations and few corrected ones, and false feedback leads the same models to delete correct quotations.

10. Ethics and Release

All source material is public court data; no generated draft was filed. Instruments, drafts, per-citation records, and the analysis ledger are at github.com/hardy30894/citation-under-pressure. Language-model assistance was used in writing code and prose under the author’s direction; every reported number is produced by the released instruments.

References

- [1] *Mata v. Avianca, Inc.*, 678 F. Supp. 3d 443 (S.D.N.Y. 2023); 2023. Opinion and order on sanctions, No. 1:22-cv-01461 (PKC), June 22, 2023.
- [2] Charlotin D. AI Hallucination Cases Database; 2025. Accessed 2 September 2026; 2,008 cases listed. <https://www.damiencharlotin.com/hallucinations/>.
- [3] Liu P, Stambach D, Henderson P. Who Checks the Citations? Benchmarking Legal Hallucination Detection; 2026. Accepted at the Conference on Language Modeling (COLM) 2026. Available from: <https://arxiv.org/abs/2606.21155>. doi:10.48550/arXiv.2606.21155.
- [4] Chen S, Yin H, Zhou S. LegalCiteBench: Evaluating Citation Reliability in Legal Language Models; 2026. Available from: <https://arxiv.org/abs/2605.10186>. doi:10.48550/arXiv.2605.10186.
- [5] Zhao C, Tang Y, Qian Y. Do Deployment Constraints Make LLMs Hallucinate Citations? An Empirical Study across Four Models and Five Prompting Regimes; 2026. Available from: <https://arxiv.org/abs/2603.07287>. doi:10.48550/arXiv.2603.07287.
- [6] Harvard Law School Library Innovation Lab. Caselaw Access Project; 2024. Digitized collection of official, book-published U.S. case law; accessed September 2026. <https://case.law/>.
- [7] Dahl M, Magesh V, Suzgun M, Ho DE. Large Legal Fictions: Profiling Legal Hallucinations in Large Language Models. *Journal of Legal Analysis*. 2024;16(1):64-93. doi:10.1093/jla/laae003.
- [8] Magesh V, Surani F, Dahl M, Suzgun M, Manning CD, Ho DE. Hallucination-Free? Assessing the Reliability of Leading AI Legal Research Tools. *Journal of Empirical Legal Studies*. 2025;22(2):216-42. doi:10.1111/jels.12413.
- [9] Gehring J, Zheng K, Copet J, Mella V, Cohen T, Synnaeve G. RLEF: Grounding Code LLMs in Execution Feedback with Reinforcement Learning. In: *Proceedings of the 42nd International Conference on Machine Learning*. vol. 267 of *Proceedings of Machine Learning Research*. PMLR; 2025. p. 19034-55. Available from: <https://proceedings.mlr.press/v267/gehring25a.html>.
- [10] Blair-Stanek A, Van Durme B. LLMs Provide Unstable Answers to Legal Questions. In: *Proceedings of the Twentieth International Conference on Artificial Intelligence and Law (ICAIL 2025)*. Chicago, IL, USA: ACM; 2025. p. 425-9. doi:10.1145/3769126.3769245.
- [11] Gultekin S, Rossi Reich M, Galloni F, Lagioia F, Consiglio E, Sartor G, et al. A Two-Dimensional Evaluation Framework for Factual and Reasoning Assessment of LLMs in Legal Question Answering. In: *Legal Knowledge and Information Systems: JURIX 2025*. vol. 416 of *Frontiers in Artificial Intelligence and Applications*. IOS Press; 2025. p. 86-97. doi:10.3233/FAIA251579.
- [12] Sanchi M, Novotná T. Towards Hybrid Evaluation Methodologies for Large Language Models in the Legal Domain. In: *Legal Knowledge and Information Systems: JURIX 2024*. vol. 395 of *Frontiers in Artificial Intelligence and Applications*. IOS Press; 2024. p. 389-92. doi:10.3233/FAIA241279.
- [13] Garneau N, Olsen HP, Corduant A, Tarissan F. Combining Network and Text to Provide Legal Pincites. In: *Legal Knowledge and Information Systems: JURIX 2024*. vol. 395 of *Frontiers in Artificial Intelligence and Applications*. IOS Press; 2024. p. 282-7. doi:10.3233/FAIA241256.
- [14] Zhang L, Savelka J, Ashley KD. Do LLMs Truly “Understand” when a Precedent Is Overruled? In: *Legal Knowledge and Information Systems: JURIX 2025*. vol. 416 of *Frontiers in Artificial Intelligence and Applications*. IOS Press; 2025. p. 239-50. doi:10.3233/FAIA251592.
- [15] Nguyen HT, Satoh K. ConsRAG: Minimize LLM Hallucinations in the Legal Domain. In: *Legal Knowledge and Information Systems: JURIX 2024*. vol. 395 of *Frontiers in Artificial Intelligence and Applications*. IOS Press; 2024. p. 327-32. doi:10.3233/FAIA241263.
- [16] Trozze A, Davies T, Kleinberg B. Large language models in cryptocurrency securities cases: can a GPT model meaningfully assist lawyers? *Artificial Intelligence and Law*. 2025;33(3):691-737. doi:10.1007/s10506-024-09399-6.
- [17] Harasta J, Novotná T, Savelka J. It cannot be right if it was written by AI: on lawyers’ preferences of documents perceived as authored by an LLM vs a human. *Artificial Intelligence and Law*. 2026;34(1):153-90. doi:10.1007/s10506-024-09422-w.
- [18] Gray MA, Savelka J, Oliver WM, Ashley KD. Generating Legal Arguments with Automatically Identified Factor Magnitudes. In: *Proceedings of the Twentieth International Conference on Artificial Intelligence and Law (ICAIL 2025)*. ACM; 2025. p. 199-208. doi:10.1145/3769126.3769212.
- [19] Pires R, Malaquias Junior R, Nogueira R. Automatic Legal Writing Evaluation of LLMs. In: *Proceedings of the Twentieth International Conference on Artificial Intelligence and Law (ICAIL 2025)*. ACM; 2025. p. 420-4. doi:10.1145/3769126.3769227.

- [20] Khatri M, Sheik R, Wadhwa P, Satija G, Kumar Y, Shah RR, et al. CiteCaseLAW: Citation Worthiness Detection in Caselaw for Legal Assistive Writing. In: Legal Knowledge and Information Systems: JURIX 2023, vol. 379 of Frontiers in Artificial Intelligence and Applications. IOS Press; 2023. p. 257-62. doi:10.3233/FAIA230971.
- [21] Sadeghian A, Sundaram L, Wang DZ, Hamilton WF, Branting K, Pfeifer C. Automatic semantic edge labeling over legal citation graphs. Artificial Intelligence and Law. 2018;26(2):127-44. doi:10.1007/s10506-018-9217-1.
- [22] Dadgostari F, Guim M, Beling PA, Livermore MA, Rockmore DN. Modeling law search as prediction. Artificial Intelligence and Law. 2021;29(1):3-34. doi:10.1007/s10506-020-09261-5.
- [23] Mumford J, Florimonte F, Atkinson K, Dzehtsiarou K. Context-Aware Citation Networks: A Human-AI Dataset, Analysis, and Tool. In: Legal Knowledge and Information Systems: JURIX 2025, vol. 416 of Frontiers in Artificial Intelligence and Applications. IOS Press; 2025. p. 145-55. doi:10.3233/FAIA251584.
- [24] Spaeth HJ, Epstein L, Nelson MJ, Martin AD, Segal JA, Ruger TJ, et al.. Supreme Court Database, Version 2025 Release 01; 2025. Modern database, terms 1946 to 2024. <https://scdb.la.psu.edu/>.
- [25] Ortega JE, Joshi DD, Borkowski MP. Large-Language Memorization During the Classification of United States Supreme Court Cases; 2025. Available from: <https://arxiv.org/abs/2512.13654>. doi:10.48550/arXiv.2512.13654.
- [26] Cushman J, Dahl M, Lissner M. eyecite: A tool for parsing legal citations. Journal of Open Source Software. 2021;6(66):3617. Available from: <https://doi.org/10.21105/joss.03617>. doi:10.21105/joss.03617.
- [27] Liang KY, Zeger SL. Longitudinal data analysis using generalized linear models. Biometrika. 1986;73(1):13-22. doi:10.1093/biomet/73.1.13.
- [28] Norman JD, Rivera MU, Hughes DA. Reliability without Validity: A Systematic, Large-Scale Evaluation of LLM-as-a-Judge Models Across Agreement, Consistency, and Bias; 2026. Available from: <https://arxiv.org/abs/2606.19544>. doi:10.48550/arXiv.2606.19544.
- [29] Zheng L, Chiang WL, Sheng Y, Zhuang S, Wu Z, Zhuang Y, et al.. Judging LLM-as-a-Judge with MT-Bench and Chatbot Arena; 2023. NeurIPS 2023 Datasets and Benchmarks Track. Available from: <https://arxiv.org/abs/2306.05685>. doi:10.48550/arXiv.2306.05685.