

Citation under Pressure: What Deployment Constraints Do to Legal Citation Integrity, and What Deterministic Verification Can Repair

Harsh DUBEY¹

New York University, New York, USA

Abstract. Sanctions over fabricated legal citations have made language-model drafting a problem for courts, and most research on the problem has produced detection benchmarks that score errors in finished briefs. This paper measures how the errors are produced. Seven models from seven vendors, from a 30B open-weight model to frontier systems, drafted argument sections for 48 U.S. Supreme Court matters under five prompting conditions: baseline, a citation quota, a restriction to pre-1970 authority, a sanctions warning, and all three combined. Every citation in the 1,680 drafts was checked deterministically, with no language-model judge in any primary measurement: existence against an 18-million-row reporter index, quotations against a verbatim checker validated on seeded faults, and pinpoint pages against official reporter pagination from public archives. The effects are specific to the model and the constraint, and the failure mode moves with capability: under combined pressure the smallest model’s citation existence fell from 0.924 to 0.773, verbatim quotation accuracy fell for four models in the pre-registered run and for six of seven once three generations are pooled, and Sonnet 5’s rate did not fall because it quoted less under pressure; its residual errors are paraphrase presented as quotation, real language from the wrong case, and wrong pages. In a second experiment the checker fed its findings back into revision. Sonnet 5 reached 1.000 verbatim accuracy, and a per-quotation trace shows how: of its 36 flagged quotations, 3 were corrected and 31 deleted. A flag that carries the closest passage of the cited opinion raises the corrections fourfold, and removal still exceeds repair in every model. False feedback produced the same deletions aimed at correct quotations, and a verifier with 50 percent precision removed 20 of Sonnet 5’s 68 correct quotations against 3 under a precise one. A second task on 48 federal appellate matters shows lower baseline existence for six of seven models. We release the instruments, including a page-level pinpoint verifier from public data.

Keywords. legal citation, hallucination, large language models, verification, evaluation

1. Introduction

Since 2023, courts have repeatedly sanctioned lawyers who filed briefs with citations a language model invented, from *Mata v. Avianca* [1] to a docket a public tracker now

¹Corresponding Author: Harsh Dubey, New York University; hd2225@nyu.edu.

counts in the thousands [2]. Most published work frames the problem as detection: given a brief that already contains errors, how many can a model or agent find [3,4]? Detection accuracy is the right measure for reviewing a filed brief; it does not tell a deployer which prompting conditions produce the errors, this paper’s question.

This paper asks two production-side questions. First, which prompting conditions cause citation failure, and for which models? Deployed prompts carry constraints of the kind Zhao et al. [5] studied for scholarly citations, where a publication-year window lowered citation existence even in capable models while a breadth demand raised it for the strongest model and lowered it for the open-weight ones; we add a sanctions warning and know of no comparable study for law. Second, when a deterministic checker feeds its findings back into revision, does the model correct them, leave them, or change the draft so that the tests pass without the quotations being corrected?

Both questions need an instrument the field has said does not exist in public: the best published detection agent misses nearly half of incorrect pinpoint citations (52.8 percent recall), its authors concluding that “no models can reliably detect wrong pincites” because page-level information lives behind commercial databases [3]. We built it from public sources: existence against a local index derived from CourtListener bulk data, verbatim quotation checking validated on seeded faults and audited on real drafts, and page-level pinpoint verification for U.S. Reports authority against the official pagination in the Caselaw Access Project’s volume archives [6].

Our primary setting is the closed-book use of general-purpose models, the setting of most sanctioned filings documented to date [2]; a grounded arm with ten retrieved, verified authorities in the prompt measures what retrieval changes. Drafts under the four non-baseline conditions are the pressured drafts (1,344 in the Supreme Court task, three refusals), and the models with the highest baseline existence are the top of the range (Section 3). The paper contributes three things: a constraint-sensitive evaluation, in which reliability at a neutral prompt does not predict reliability under drafting constraints and the failure mode moves with capability; deterministic verification of existence, verbatim quotation, and pinpoint page with no language-model judge, released with the drafts; and the finding that verification is not repair, since flags fed back to the model raise measured accuracy by deletion and an imprecise verifier damages correct content.

2. Related Work

Legal citation reliability. Liu et al. [3] build a detection benchmark from real appellate briefs with injected errors and a taxonomy grounded in sanctioned filings; we adopt their error categories and measure the naturally occurring errors their limitations section asks for. LegalCiteBench [4] shows that models asked closed-book for a named case’s citation almost never produce it or abstain; our task asks the easier question, whether the cases a model chooses to cite exist. Dahl et al. [7] and Magesh et al. [8] measure hallucinated authority in chatbots and research tools.

Constraints and fabrication. Zhao et al. [5] is the closest design: five prompting regimes over scholarly claims, verified by fuzzy bibliographic matching that left 36 to 61 percent of citations unresolved. We transplant their factorial to law with deterministic verification (0.2 percent unresolved) and a formal statistical model.

Revision with execution feedback. The loop protocol follows execution-feedback work in code generation [9], with its no-feedback control separating feedback content from extra revision; we add controls in which a quarter to all of the feedback is false.

AI and law. Blair-Stanek and Van Durme [10] found at ICAIL that frontier models flip the winning party across identical zero-temperature runs, which Section 8 bounds with repeated generations. Evaluation in this community pairs automated metrics with the checks lawyers perform [11,12]; citation integrity is one with deterministic ground truth. Garneau et al. [13] found retrieving the paragraph a proposition rests on far harder than retrieving the case; our verifier meets the same gap from the other side, checking a generated pincite against the page text. Zhang et al. [14] show models miss overruled precedent; ConsRAG [15] constrains retrieval, complementary to our loop’s feedback; Trozze et al. [16] and Harasta et al. [17] study how jurors and lawyers perceive model-written documents, a perception machine verification does not depend on. Evaluations of generated legal argument [18,19] do not check the authority it cites, the layer this paper adds; citation-worthiness and citation-network work [20,21,22,23] studies the network our checker queries.

3. Tasks, Conditions, and Instruments

Terms. In 347 U.S. 483, 495, the volume, reporter, and first page identify the case and the second page is the *pincite*, for a quotation the page where the quoted language appears. *Citation existence* is the share of adjudicable citations that resolve to a real case; *strict quotation accuracy* the share of scored quotations, those the draft attributes to a case whose text the checker holds, that reproduce the opinion verbatim; a quotation the draft does not attribute is outside the rate.

Matters. The matters are forty-eight Supreme Court cases from 1990 to 2020 drawn from the Supreme Court Database [24]. Each packet holds the question presented, party names, a factual summary, and a lower-court opinion excerpt, the Supreme Court’s decision excluded and outcome leakage screened. The model argues an assigned side (petitioner for even-numbered matters, respondent for odd) in an argument section of under 900 words; nothing predicts outcomes, so the memorization confound documented for Supreme Court case classification [25] is out of scope. A second task uses forty-eight published F.3d decisions from the Caselaw Access Project archive (1995 to 2019, four from each of the First through Eleventh and Federal Circuits, no cross-appeals). The packet is the opinion’s own statement of the case and facts, cut where the analysis begins and scrubbed of outcome sentences, and the model argues for the named appellant under all five conditions.

Conditions. The five conditions are baseline; quota (“cite at least eight distinct case authorities”); temporal (“every authority you cite MUST have been decided before 1970”), a clause no user writes that stands in for the binding-authority and jurisdiction limits practitioners do impose, and which, unlike Zhao et al.’s window on recent work, moves citations toward well-memorized opinions; stakes (“this section will be filed with the court; citing authority that does not exist, or misquoting authority, risks sanctions under Rule 11”); and combined, which concatenates the three clauses. Each condition has one fixed template in the primary run, at temperature 0. The stakes clause names Rule 11, the rule invoked in the sanctioned-filing cases [1], which governs district-court

civil actions [26], so for a Supreme Court filing the clause is a stakes signal only. Each clause is a single sentence so that its effect can be isolated; a second template rephrasing every condition was run on all 48 matters and seven models. A *grounded* arm reran baseline and combined with ten retrieved authorities in the prompt: the U.S. Reports opinions most similar to the question and facts (TF-IDF over the first 2,500 words of the 25,250 cached opinions of 2,000 characters or more with a name and date in the index, decided before argument and before 1970 for the combined condition, neither party named), each with citation, year, and its best-matching 250-word passage, stated to be real, quotation restricted to the excerpts.

Models. The models are Qwen3-30B-A3B, Mistral Small, Llama-4 Maverick, DeepSeek V4 Flash, Grok 4.3, GPT-5.4-mini, and Claude Sonnet 5, one per vendor, spanning open-weight 30B to frontier; tables are ordered by baseline strict quotation accuracy, and Sonnet 5, GPT-5.4-mini, and Grok 4.3, with the highest baseline existence, are called the top of the range. The seven were run through OpenRouter in August 2026 with completion limits of 1,400 to 12,000 tokens; the identifiers are qwen3-30b-a3b-instruct-2507, mistral-small-2603, llama-4-maverick, deepseek-v4-flash, grok-4.3, gpt-5.4-mini, and claude-sonnet-5. An eighth model, GLM-4.7-flash, produced no output on 33 of 240 drafts (28 combined), exhausting a 30,000-token budget in reasoning, a third failure mode beside fabrication and refusal; it is excluded from the tables.

Adjudication. Citations are extracted with eyecite [27] and given one of three labels: exists, not found, or unverifiable (vendor citations and coverage gaps), so that index coverage is never counted as fabrication; 0.2 percent of the 14,040 citation records in the full run are unverifiable. Citations resolve against an 18-million-row table; quotations are located by balanced quotation marks and attributed, in order, to the citation whose parenthetical holds them, the case named in their own sentence, the citation that follows them in that sentence (not one introduced by *see* or *cf.*), the citation an *Id.* points to, or the last citation in the paragraph; a quotation introduced by a statute, rule, the record, or a lower court, or with no citation in its paragraph, is left unpaired, since the checker has nothing to test it against: 37 percent of the 10,846 quotations in the full run, 24 percent introduced by non-case material and 10 percent standing in a paragraph that cites nothing, a share that runs from 31 to 42 percent across models. Quotations are checked verbatim against the cited opinion’s text from public sources. *Exists* means the volume, reporter, and page resolve to an indexed opinion; the case name is compared separately, and 205 of 11,769 resolving citations (1.7 percent) carry a name the index does not match, a ceiling that includes abbreviated names (Sonnet 5: 6 of 1,787). A quotation is *near miss* when every fragment reaches token coverage of at least 0.85 without matching exactly, and *inaccurate* otherwise. The strict rate for a draft d is $S_d = n_{\text{acc}} / (n_{\text{acc}} + n_{\text{near}} + n_{\text{inacc}})$, and a lenient rate counting near misses as correct is reported alongside it. Matching is case and punctuation insensitive, splits on ellipses and on bracketed alterations (so that a lawful substitution is read as an omission and the words around it must match), and drops alteration parentheticals; the lenient rate is a diagnostic, since 0.85 coverage admits one wrong word in seven. The quotation checker scores 240 of 240 on twenty synthetic drafts planting twelve item kinds, six of them attribution traps (a rival citation after a signal phrase, a statute, an *Id.*, a string-cite parenthetical, paragraphs with and without a citation), and the existence and pinpoint checks 100 of 100 on 20 seeded opinions; Section 6 audits the checker on real drafts. Pinpoint pages are verified against the official




Figure 1. Strict quotation accuracy (left) and citation existence (right) by condition, one line per model; the dotted line is the human-lawyer floor under the same instrument (Section 3); the right axis starts at 0.7.

reporter pagination in the Caselaw Access Project’s HTML volumes, which mark every page break; a quotation counts as on the cited page when its longest fragment lies within that page’s text.

Human reference. The same instrument scored 482 pre-ChatGPT appellate brief excerpts from Liu et al. [3] (no injected error; OCR-converted, so a floor); human lawyers reach 0.594 strict and 0.776 lenient quotation accuracy (416 scored quotations) and 0.957 citation existence (0.976 on U.S. Reports and published federal reporters; Federal Appendix spellings, once mismatched against the index, now resolve for models and briefs alike), a reference for the severity of the verbatim standard only, since human briefs differ in citation mix. The literal matcher the strict standard replaced, which read a bracketed alteration as its contents, puts the human rate at 0.493 and moves model cells by at most 9 points, with the same 8 contrasts surviving correction (Section 4); of the 169 human strict failures that remain, 59 are near misses without any alteration, the signature of a corrupted character.

Inference. The primary analysis is a citation-level logistic generalized estimating equation per model, $\text{logit } P(y_{imk} = 1) = \beta_0 + \beta_k$ for citation i in matter m under condition k , with exchangeable working correlation within matters [28], so that the 48 matters are the independent units. Each model has 8 contrasts (four conditions, two outcomes); we report Holm-corrected p values within model, each model being the family the paper draws conclusions about, and call an effect significant only when it survives correction; counts of models (“three of seven”) are descriptive, and with 48 matters as the unit the effective sample is 48, not the thousands of citations. Odds ratios, intervals, and p values for all 56 contrasts are released, and a matter-level cluster bootstrap agreed in sign with every corrected result. Three robustness fits follow the same model: one pools two further generations of the baseline and combined drafts (Section 8) with a run effect (Holm over 14 tests), the others pool both templates or both tasks with a template or task effect, corrected within model. For the revision experiment, pre-registered descriptively, we add matter-paired comparisons of final strict accuracy between arms (Wilcoxon, bootstrap interval). Hypotheses and decision rules were frozen before the full run; every later decision, including two withdrawn findings, is in a dated ledger.

4. Experiment 1: Pressure

Table 1 and Figure 1 report the rates.

Citation existence. In the single run, only the 30B model’s existence moves significantly under pressure: from 0.924 at baseline to 0.773 under the combined condition,

Table 1. Baseline versus the combined condition; quotation columns give the strict rate over attributed quotations (lenient in parentheses) [scored quotations], existence columns the rate [citations not found / adjudicable], the last column the share of checked citations decided after 1969 under the temporal condition; an asterisk marks a contrast significant after Holm correction within model.

Model	Quotes, base	Quotes, comb.	Exist, base	Exist, comb.	Viol., temporal
Qwen3-30B	0.165 (0.609) [230]	0.061 (0.320) [325]*	0.924 [28/370]	0.773 [90/397]*	20%
Mistral	0.262 (0.620) [187]	0.157 (0.444) [248]	0.901 [37/373]	0.868 [69/521]	35%
DeepSeek	0.381 (0.730) [252]	0.142 (0.491) [281]*	0.972 [10/354]	0.959 [20/488]	11%
Grok 4.3	0.455 (0.737) [156]	0.302 (0.647) [119]	0.981 [5/257]	0.994 [3/479]	0%
Sonnet 5	0.523 (0.804) [260]	0.644 (0.825) [194]	0.991 [3/318]	0.998 [1/450]	0%
Llama-4	0.588 (0.784) [51]	0.230 (0.557) [61]*	0.958 [9/213]	0.939 [31/506]	10%
GPT-5.4-mini	0.583 (0.819) [127]	0.466 (0.693) [88]	0.981 [8/417]	0.962 [27/713]	13%

with the quota, temporal, and combined contrasts all significant after correction (combined OR 0.30, 95 percent CI 0.18 to 0.51, corrected $p < 0.001$). Mistral Small falls under the temporal and combined conditions without reaching significance, and the five other models stay between 0.93 and 1.00 in every condition. Each surviving existence contrast rests on at least 28 citations not found per arm; a record is one occurrence of a full citation (short forms are not counted), and refitting on distinct authorities per draft leaves the same 8 contrasts.

Quotation accuracy. The temporal restriction damages quoting most; after correction, strict quotation accuracy falls under temporal and combined pressure for Llama-4 (0.588 to 0.230), under the combined condition for Qwen3-30B (0.165 to 0.061) and DeepSeek (0.381 to 0.142), and under the temporal clause for Mistral Small; GPT-5.4-mini and Grok 4.3 fall in the same direction without surviving correction in the single run, and Sonnet 5 does not fall (below). Llama-4 shows the largest drop. At baseline its strict rate is 0.588 (the human floor is 0.594), and under the temporal restriction it drops to 0.214 (OR 0.18, CI 0.08 to 0.43, corrected $p < 0.001$); refitting every contrast on the lenient outcome, 8 survive correction, the combined fall among them for Qwen3-30B, DeepSeek, and Mistral Small. Scored quotations per cell run from 42 to 325 and the unpaired share from 15 to 50 percent, so the rates are conditional on quoting, the point Sonnet 5’s contrast turns on (below). Nor is the fall an artifact of scanned text for old opinions; at baseline, strict accuracy for quotations from opinions decided before 1970 is 0.439 (189 quotations) against 0.383 for later ones (1,074). With three generations pooled, the combined-condition fall survives correction for 6 of the seven models (OR 0.20 to 0.53), Sonnet 5 excepted, and Mistral Small’s existence fall also survives (OR 0.62, corrected $p=0.004$).

Robustness to phrasing. Under the second template, 12 contrasts survive correction, among them 5 of the 8 that survived under the first (Qwen3-30B’s quota existence fall and Llama-4’s two quotation falls did not); pooling both templates, all 8 survive and the combined-condition fall in quotation accuracy survives for 6 of seven models (OR 0.27 to 0.46), Sonnet 5 again the exception (pooled OR 1.28, corrected $p=0.963$). Cell rates differ between templates by up to 15 points of quotation accuracy and 19 of existence.

Second task: federal appellate briefs. Outside the Supreme Court, fabrication at a neutral prompt is more common. Baseline citation existence on the 48 appellate matters runs from 0.652 for Qwen3-30B to 0.993 for Sonnet 5, lower than on the Supreme Court

matters for 6 of the seven models, Sonnet 5 alone unchanged. The failures concentrate in the federal reporters: 22 percent of citations to F.2d, F.3d, and F. Supp. do not resolve, against 4 percent for the U.S. Reports; the index resolves all but 3 of the 759 federal-reporter citations the 48 deciding opinions themselves contain, so the shortfall is not coverage. One pressure contrast survives correction within the task, Llama-4’s existence fall under the temporal clause (0.950 to 0.857, OR 0.32, corrected $p=0.018$). Pooling both tasks with a task effect, seven of the 8 Supreme Court survivors remain (Qwen3-30B’s quota existence fall drops out), joined by Llama-4’s existence falls under temporal and combined, DeepSeek’s temporal and Mistral Small’s combined quotation falls, and Sonnet 5’s rate rise under the stakes clause (OR 1.53, corrected $p=0.020$).

Grounded drafting. With ten verified authorities in the prompt, the models drew 51 to 83 percent of their citations from the list; citations not found fell from 341 in 5,856 closed-book to 74 in 5,605, and Qwen3-30B’s combined-condition existence went from 0.773 to 0.995 (OR 53.8). Quotation accuracy rose in every cell (Qwen3-30B combined 0.061 to 0.412; Sonnet 5 baseline 0.523 to 0.712), yet with the source passages in the prompt no cell exceeded 0.734, and 61 percent of Sonnet 5’s 71 remaining inaccurate quotations are paraphrases of the cited case in quotation marks.

The strongest model does not degrade, because it quotes less. Sonnet 5’s strict rate is numerically higher under the temporal, stakes, and combined conditions and highest under the combined one (0.523 to 0.644; uncorrected $p = 0.0313$, corrected $p = 0.250$), the one pre-registered contrast in that direction, but the rate is conditional on quoting: Sonnet 5 wrote 9.6 quotations per draft at baseline and 6.1 under the combined condition, and its accurate quotations per draft fell from 2.83 to 2.60. The rise is selection, fewer and safer quotations, and the temporal restriction also shifts the cited cases toward older, better-memorized opinions (median decision year 1943 against 1986). With three generations pooled Sonnet 5’s existence rise survives the logistic fit (corrected $p=0.006$), but it rests on twelve citations not found (7 of 975 at baseline, 5 of 1,324 under combined), on which an exact test gives $p=0.38$; the quotation rise reaches corrected $p=0.404$ there. In the appellate task the stakes clause alone raises the rate (0.480 to 0.620, corrected $p=0.084$) without moving the cited cases (median decision year 1991 at baseline and under stakes); the direction is consistent across tasks and its cause, under selection, is open.

5. Experiment 2: The Checker inside the Revision Loop

Combined-condition drafts entered a revision loop: each round the checker returned one line per failure (a citation that does not resolve; a quotation not verbatim in the cited opinion), the model revised, and the loop ended on a clean draft or after three revisions. A *false feedback* control gave the same number of identically formatted lines aimed at correct items, *mixed* controls replaced each line with a false one at random for a precision of about 0.25, 0.5, or 0.75, and a *no-feedback* control gave the same rounds with a generic instruction to revise. The loop used the first 24 matters in pre-registered order; loop rates are means over drafts with at least one scored quotation (a draft left with none drops out; 68 of the 168 true-arm episodes ended so, 19 of Grok 4.3’s 24, so its final rate rests on 5 drafts), so round-0 figures differ from the pooled rates in Table 1. The loop was rerun after the attribution rule was corrected; the original run, whose true-arm flags had

Table 2. Revision under true checker feedback, 24 episodes per model. Round-0 quotations, flagged (inaccurate or near miss) and accurate, traced to the final draft; corr.: corrected, an accurate quotation on the same citation now stands in its place, edited or new; left: still present and flagged; removed: deleted, dequoted, or made inaccurate; added: quotations new in the final draft, accurate/flagged; strict accuracy at round 0 and final as the draft mean (pooled over quotations); Clean: episodes ending on a passing draft.

Model	Flagged at round 0					Accurate at round 0					Clean
	n	corr.	dequoted	deleted	left	n	removed	added	Strict, r0 → final (pooled)		
Qwen3-30B	138	0	10	13	115	12	0	0/32	0.089 → 0.092 (0.075)	0	
Mistral	96	4	4	43	45	13	1	1/27	0.233 → 0.350 (0.191)	10	
DeepSeek	136	2	59	71	4	29	4	1/2	0.292 → 0.950 (0.824)	23	
Grok 4.3	37	0	23	14	0	13	3	0/0	0.234 → 1.000 (1.000)	24	
Sonnet 5	36	3	2	31	0	68	3	0/0	0.737 → 1.000 (1.000)	24	
Llama-4	33	0	5	28	0	11	2	0/0	0.273 → 1.000 (1.000)	24	
GPT-5.4-mini	22	1	1	19	1	31	5	0/1	0.538 → 0.857 (0.931)	20	




Figure 2. Verifier precision against what revision does. Left: final strict quotation accuracy. Right: share of accurate round-0 quotations that revision deleted, dequoted, or made inaccurate. Precision is the share of true feedback lines; the leftmost point is revision with no feedback.

realized precision 0.71 under the corrected rule (0.38 for Sonnet 5), is released beside it. Every scored round-0 quotation was traced to the final draft as corrected, dequoted, deleted, or left in place (Table 2).

Satisfying the checker by deletion. Under true feedback the final strict rate ranges from 0.092 for Qwen3-30B to 1.000 for Sonnet 5 (Table 2). Sonnet 5 had 36 flagged quotations at round 0: 1 was corrected in place, 2 replaced by a new accurate quotation from the same case, 2 dequoted, and 31 deleted, while only 3 of its 68 accurate quotations went, against 21 under revision with no feedback. Every model but Qwen3-30B and Mistral Small, which left most flags standing, removed most of what was flagged. Citation existence rose by the same route: of 113 citations unresolved at round 0, 82 were gone from the final draft, 31 remained, and 37 new ones resolved. Rerun from the grounded combined drafts with the retrieved excerpts in the prompt, the loop repaired 4 of 23 flags for Sonnet 5 and 1 of 51 for DeepSeek. A *passage* arm then put the evidence in the flag itself: each quotation line carried the closest passage of the cited opinion (Section 3’s window search). Repair rose in six of seven models (Grok 4.3 still repaired none): to 12 of 36 for Sonnet 5 (3 in place, 9 by the supplied passage), 26 of 96 for Mistral Small against 4, 24 of 138 for Qwen3-30B against none; removal still exceeded it in every model (23 for Sonnet 5). A loop should hand the model the passage; even then it removes more than it repairs.

Controls: feedback content and verifier precision. Without feedback (Figure 2), Sonnet 5’s strict rate still rose (0.737 to 0.787), and true feedback added 0.21 over revision alone (paired over matters scored in both arms; interval 0.12 to 0.32, $p=0.002$), as for DeepSeek (0.38); GPT-5.4-mini’s, Qwen3-30B’s, and Mistral Small’s differences (0.04, 0.02, 0.05) are within noise; Llama-4 and Grok 4.3 left too few paired matters. Sonnet 5’s correct quotations removed rise as precision falls, 3, 8, 20, 27, and 43 of 68 at precisions of 1, 0.75, 0.5, 0.25, and 0, against 21 with no feedback. The mixed arms lower recall with precision; the right panel isolates the false-flag effect. Under all-false feedback five of seven models ended below their starting rate while flagged counts fell in every model: what changed was the deletion of correct quotations.

6. Residual Errors at the Top of the Range

Paraphrase inside quotation marks. Decomposing every inaccurate quotation against the cited opinion’s text, 68 percent of Sonnet 5’s (169 of 250) are paraphrases of the correct case presented as verbatim quotation, 56 match nothing substantial in the cited opinion, and 25 cite authority outside the U.S. Reports, beyond the decomposition’s corpus.

Real language, wrong case. Searching every inaccurate quotation against the 62,049 U.S. Reports opinions in the Caselaw Access Project cache locates 326 quotations across the seven models (72 for Sonnet 5, 43 for Qwen3-30B) whose language appears verbatim in a different opinion than the one cited (in 139 the true source is cited in the same draft, so the count is a ceiling); the rest, found nowhere or matching only stock phrases (130 and 43 for Sonnet 5), are the paraphrases above and quotations from other courts. Two audits bound the checker’s error: the same search finds the language present in the attributed case for 5 of Sonnet 5’s 250, and a hand reading of 160 inaccurate verdicts drawn at random across models found 35 on the checker’s side (22 percent, 95 percent interval 16 to 29: non-case material in quotation marks, or a quotation bound to the wrong case), 111 the model’s, 4 failing only on a bracketed alteration, and 10 open. The share was lowest under the combined condition (4 of 42) and highest under the quota (11 of 30), so residual error works against the falls reported; 40 accurate verdicts read the same way all held. The rule’s first version put about half of a forty-item sample on the checker’s side.

Right case, wrong page. Of the 9,129 pinpoints in the Supreme Court drafts, 73 percent were checked against the cited case’s page span (the rest cite other reporters or cases past volume 572, the archive’s last) and 1,006 carry a quotation on the same citation; for those the quotation lies on the cited page 83 percent of the time for Sonnet 5, 50 to 67 percent for the other six, and 88 percent for the human briefs; 12 percent lie on an adjacent page and the rest elsewhere in the cited case.

7. Discussion

Where the checker belongs. The revision experiment favors the checker at the gate, returning failures to a person, over a loop in which the model resolves bare flags, and says what a loop should carry if built, the passage; removing an unverifiable quotation

is a defensible response to a flag, but the quotation is deleted or survives as untestable paraphrase, false flags get the same response, and a loop-pattern tool reports a clean draft.

Verifier precision is a deployment requirement. Our checker’s error among inaccurate verdicts is 22 percent by the hand reading (Section 6), and published detectors are less precise [3]. Figure 2 shows what false flags do: at precision one half Sonnet 5 lost 20 of its 68 correct quotations, against 3 under a precise verifier and 21 with no feedback at all. A pipeline that feeds a probabilistic detector’s flags to a model should report the detector’s precision beside its recall, and a deployer should enforce the strict standard as defined here, which admits bracketed alterations, with the 0.85-coverage rate as a diagnostic.

What evaluations should include. The two experiments make one point. A model can raise an observable correctness rate by changing what it exposes to the check, quoting less under pressure (Section 4) and deleting under feedback (Section 5), without the underlying quoting improving; a rate should therefore be read beside the count of correct items it is computed over. A single fabrication rate at a neutral prompt also misses existence under constraints [5] and on other courts, quotation accuracy, and the residual classes at the top of the range; the checks run on any filing [1,2].

8. Limitations

Both tasks are United States federal case law; the stakes clause names a sanctions rule that does not govern the forum; and the checks stop at existence, wording, and page; three model panels for relevance, calibrated on expert labels [3], fell short of a 0.75 bar (0.500 to 0.700), so none was used [29,30]. The controls, the grounded arm and loop, the second task, the per-quotation trace, and the corrected attribution rule (Section 6) all postdate the pre-registered analysis; their tests were chosen after the data were seen, every contrast was recomputed under the corrected rule, and they are reported apart from the pre-registered fit. Three generations of the baseline and combined cells (672 drafts per run) bound run-to-run variation [10]: cell standard deviations were at most 0.050 (existence) and 0.057 (strict), and the combined-versus-baseline quotation contrast kept its sign in every run for all 7 models; since the between-template spread of cells (15 points of quotation accuracy) exceeds many single-run effects, we rely on the contrasts that survive correction across runs and templates. The seven models are one per vendor, so capability, family, and recipe are confounded.

9. Conclusion

Under a quota, a date restriction, and a sanctions warning the failure mode moves with capability: the 30B model loses existence, the middle of the range verbatim quotation, and the strongest model’s rate holds because it quotes less, its residue paraphrase in quotation marks, misattribution, and wrong pages. Returned to the model, the checker’s findings raise existence chiefly by removal, flagged quotations are removed far more often than corrected even with the passage in the flag, and false flags remove correct ones; verification belongs at the gate, and verifier precision in the evaluation of any system that acts on it. Instruments, drafts, records, and ledger are released at <https://github.com/hardy30894/citation-under-pressure>.

References

- [1] *Mata v. Avianca, Inc.*, 678 F. Supp. 3d 443 (S.D.N.Y. 2023); 2023. Opinion and order on sanctions, No. 1:22-cv-01461 (PKC), June 22, 2023.
- [2] Charlotin D. AI Hallucination Cases Database; 2025. Accessed 2 September 2026; 2,008 cases listed. <https://www.damiencharlotin.com/hallucinations/>.
- [3] Liu P, Stambach D, Henderson P. Who Checks the Citations? Benchmarking Legal Hallucination Detection; 2026. Accepted at the Conference on Language Modeling (COLM) 2026. Available from: <https://arxiv.org/abs/2606.21155>. doi:10.48550/arXiv.2606.21155.
- [4] Chen S, Yin H, Zhou S. LegalCiteBench: Evaluating Citation Reliability in Legal Language Models; 2026. Available from: <https://arxiv.org/abs/2605.10186>. doi:10.48550/arXiv.2605.10186.
- [5] Zhao C, Tang Y, Qian Y. Do Deployment Constraints Make LLMs Hallucinate Citations? An Empirical Study across Four Models and Five Prompting Regimes; 2026. Available from: <https://arxiv.org/abs/2603.07287>. doi:10.48550/arXiv.2603.07287.
- [6] Harvard Law School Library Innovation Lab. Caselaw Access Project; 2024. Digitized collection of official, book-published U.S. case law; accessed September 2026. <https://case.law/>.
- [7] Dahl M, Magesh V, Suzgun M, Ho DE. Large Legal Fictions: Profiling Legal Hallucinations in Large Language Models. *Journal of Legal Analysis*. 2024;16(1):64-93. doi:10.1093/jla/laae003.
- [8] Magesh V, Surani F, Dahl M, Suzgun M, Manning CD, Ho DE. Hallucination-Free? Assessing the Reliability of Leading AI Legal Research Tools. *Journal of Empirical Legal Studies*. 2025;22(2):216-42. doi:10.1111/jels.12413.
- [9] Gehring J, Zheng K, Copet J, Mella V, Cohen T, Synnaeve G. RLEF: Grounding Code LLMs in Execution Feedback with Reinforcement Learning. In: *Proceedings of the 42nd International Conference on Machine Learning*. vol. 267 of *Proceedings of Machine Learning Research*. PMLR; 2025. p. 19034-55. Available from: <https://proceedings.mlr.press/v267/gehring25a.html>.
- [10] Blair-Stanek A, Van Durme B. LLMs Provide Unstable Answers to Legal Questions. In: *Proceedings of the Twentieth International Conference on Artificial Intelligence and Law (ICAIL 2025)*. Chicago, IL, USA: ACM; 2025. p. 425-9. doi:10.1145/3769126.3769245.
- [11] Gultekin S, Rossi Reich M, Galloni F, Lagioia F, Consiglio E, Sartor G, et al. A Two-Dimensional Evaluation Framework for Factual and Reasoning Assessment of LLMs in Legal Question Answering. In: *Legal Knowledge and Information Systems: JURIX 2025*. vol. 416 of *Frontiers in Artificial Intelligence and Applications*. IOS Press; 2025. p. 86-97. doi:10.3233/FAIA251579.
- [12] Sanchi M, Novotná T. Towards Hybrid Evaluation Methodologies for Large Language Models in the Legal Domain. In: *Legal Knowledge and Information Systems: JURIX 2024*. vol. 395 of *Frontiers in Artificial Intelligence and Applications*. IOS Press; 2024. p. 389-92. doi:10.3233/FAIA241279.
- [13] Garneau N, Olsen HP, Corduant A, Tarissan F. Combining Network and Text to Provide Legal Pincites. In: *Legal Knowledge and Information Systems: JURIX 2024*. vol. 395 of *Frontiers in Artificial Intelligence and Applications*. IOS Press; 2024. p. 282-7. doi:10.3233/FAIA241256.
- [14] Zhang L, Savelka J, Ashley KD. Do LLMs Truly “Understand” when a Precedent Is Overruled? In: *Legal Knowledge and Information Systems: JURIX 2025*. vol. 416 of *Frontiers in Artificial Intelligence and Applications*. IOS Press; 2025. p. 239-50. doi:10.3233/FAIA251592.
- [15] Nguyen HT, Satoh K. ConsRAG: Minimize LLM Hallucinations in the Legal Domain. In: *Legal Knowledge and Information Systems: JURIX 2024*. vol. 395 of *Frontiers in Artificial Intelligence and Applications*. IOS Press; 2024. p. 327-32. doi:10.3233/FAIA241263.
- [16] Trozze A, Davies T, Kleinberg B. Large language models in cryptocurrency securities cases: can a GPT model meaningfully assist lawyers? *Artificial Intelligence and Law*. 2025;33(3):691-737. doi:10.1007/s10506-024-09399-6.
- [17] Harasta J, Novotná T, Savelka J. It cannot be right if it was written by AI: on lawyers’ preferences of documents perceived as authored by an LLM vs a human. *Artificial Intelligence and Law*. 2026;34(1):153-90. doi:10.1007/s10506-024-09422-w.
- [18] Gray MA, Savelka J, Oliver WM, Ashley KD. Generating Legal Arguments with Automatically Identified Factor Magnitudes. In: *Proceedings of the Twentieth International Conference on Artificial Intelligence and Law (ICAIL 2025)*. ACM; 2025. p. 199-208. doi:10.1145/3769126.3769212.
- [19] Pires R, Malaquias Junior R, Nogueira R. Automatic Legal Writing Evaluation of LLMs. In: *Proceedings of the Twentieth International Conference on Artificial Intelligence and Law (ICAIL 2025)*. ACM; 2025. p. 420-4. doi:10.1145/3769126.3769227.

- [20] Khatri M, Sheik R, Wadhwa P, Satija G, Kumar Y, Shah RR, et al. CiteCaseLAW: Citation Worthiness Detection in Caselaw for Legal Assistive Writing. In: Legal Knowledge and Information Systems: JURIX 2023, vol. 379 of Frontiers in Artificial Intelligence and Applications. IOS Press; 2023. p. 257-62. doi:10.3233/FAIA230971.
- [21] Sadeghian A, Sundaram L, Wang DZ, Hamilton WF, Branting K, Pfeifer C. Automatic semantic edge labeling over legal citation graphs. Artificial Intelligence and Law. 2018;26(2):127-44. doi:10.1007/s10506-018-9217-1.
- [22] Dadgostari F, Guim M, Beling PA, Livermore MA, Rockmore DN. Modeling law search as prediction. Artificial Intelligence and Law. 2021;29(1):3-34. doi:10.1007/s10506-020-09261-5.
- [23] Mumford J, Florimonte F, Atkinson K, Dzehtsiarou K. Context-Aware Citation Networks: A Human-AI Dataset, Analysis, and Tool. In: Legal Knowledge and Information Systems: JURIX 2025, vol. 416 of Frontiers in Artificial Intelligence and Applications. IOS Press; 2025. p. 145-55. doi:10.3233/FAIA251584.
- [24] Spaeth HJ, Epstein L, Nelson MJ, Martin AD, Segal JA, Ruger TJ, et al. Supreme Court Database, Version 2025 Release 01; 2025. Modern database, terms 1946 to 2024. <https://scdb.la.psu.edu/>.
- [25] Ortega JE, Joshi DD, Borkowski MP. Large-Language Memorization During the Classification of United States Supreme Court Cases; 2025. Available from: <https://arxiv.org/abs/2512.13654>. doi:10.48550/arXiv.2512.13654.
- [26] Federal Rules of Civil Procedure, Rule 1 (Scope and Purpose); 2024. "These rules govern the procedure in all civil actions and proceedings in the United States district courts". https://www.law.cornell.edu/rules/frcp/rule_1.
- [27] Cushman J, Dahl M, Lissner M. eyecite: A tool for parsing legal citations. Journal of Open Source Software. 2021;6(66):3617. Available from: <https://doi.org/10.21105/joss.03617>. doi:10.21105/joss.03617.
- [28] Liang KY, Zeger SL. Longitudinal data analysis using generalized linear models. Biometrika. 1986;73(1):13-22. doi:10.1093/biomet/73.1.13.
- [29] Norman JD, Rivera MU, Hughes DA. Reliability without Validity: A Systematic, Large-Scale Evaluation of LLM-as-a-Judge Models Across Agreement, Consistency, and Bias; 2026. Available from: <https://arxiv.org/abs/2606.19544>. doi:10.48550/arXiv.2606.19544.
- [30] Zheng L, Chiang WL, Sheng Y, Zhuang S, Wu Z, Zhuang Y, et al. Judging LLM-as-a-Judge with MT-Bench and Chatbot Arena; 2023. NeurIPS 2023 Datasets and Benchmarks Track. Available from: <https://arxiv.org/abs/2306.05685>. doi:10.48550/arXiv.2306.05685.